

# CVA6 Hypervisor Extension Support “CVA6-H”

## Project Concept Proposal

Bruno Sá, José Martins, Sandro Pinto

Universidade do Minho / Zero-Day-Labs



OpenHW Group @ San Jose

December 15th, 2022

# Project High-Level Overview

---

# RISC-V Hypervisor Extension

- The Hypervisor Extension was ratified and frozen in Q4 2021
- Designed for type-1 and type-2 hypervisors
  - Nested virtualization is also supported
- Additional orthogonal execution modes
  - HS-mode (hypervisor-extended supervisor)
  - VS-mode & VU-mode
- Additional CSR registers:
  - HS-mode CSRs for hypervisor capabilities (e.g., hstatus)
  - HS-mode CSRs for accessing Guest/VM state (e.g., vsstatus)

## Chapter 8

### Hypervisor Extension, Version 1.0

This chapter describes the RISC-V hypervisor extension, which virtualizes the supervisor-level architecture to support the efficient hosting of guest operating systems atop a type-1 or type-2 hypervisor. The hypervisor extension changes supervisor mode into *hypervisor-extended supervisor mode* (HS-mode, or *hypervisor mode* for short), where a hypervisor or a hosting-capable operating system runs. The hypervisor extension also adds another stage of address translation, from *guest physical addresses* to supervisor physical addresses, to virtualize the memory and memory-mapped I/O subsystems for a guest operating system. HS-mode acts the same as S-mode, but with additional instructions and CSRs that control the new stage of address translation and support hosting a guest OS in virtual S-mode (VS-mode). Regular S-mode operating systems can execute without modification either in HS-mode or as VS-mode guests.

In HS-mode, an OS or hypervisor interacts with the machine through the same SBI as an OS normally does from S-mode. An HS-mode hypervisor is expected to implement the SBI for its VS-mode guest.

The hypervisor extension depends on an “T” base integer ISA with 32 *x* registers (RV32I or RV64I), not RV32E, which has only 16 *x* registers. CSR *mtval* must not be read-only zero, and standard page-based address translation must be supported, either Sv32 for RV32, or a minimum of Sv39 for RV64.

The hypervisor extension is enabled by setting bit 7 in the *misal* CSR, which corresponds to the letter H. RISC-V harts that implement the hypervisor extension are encouraged not to hardwire *misal*[7], so that the extension may be disabled.

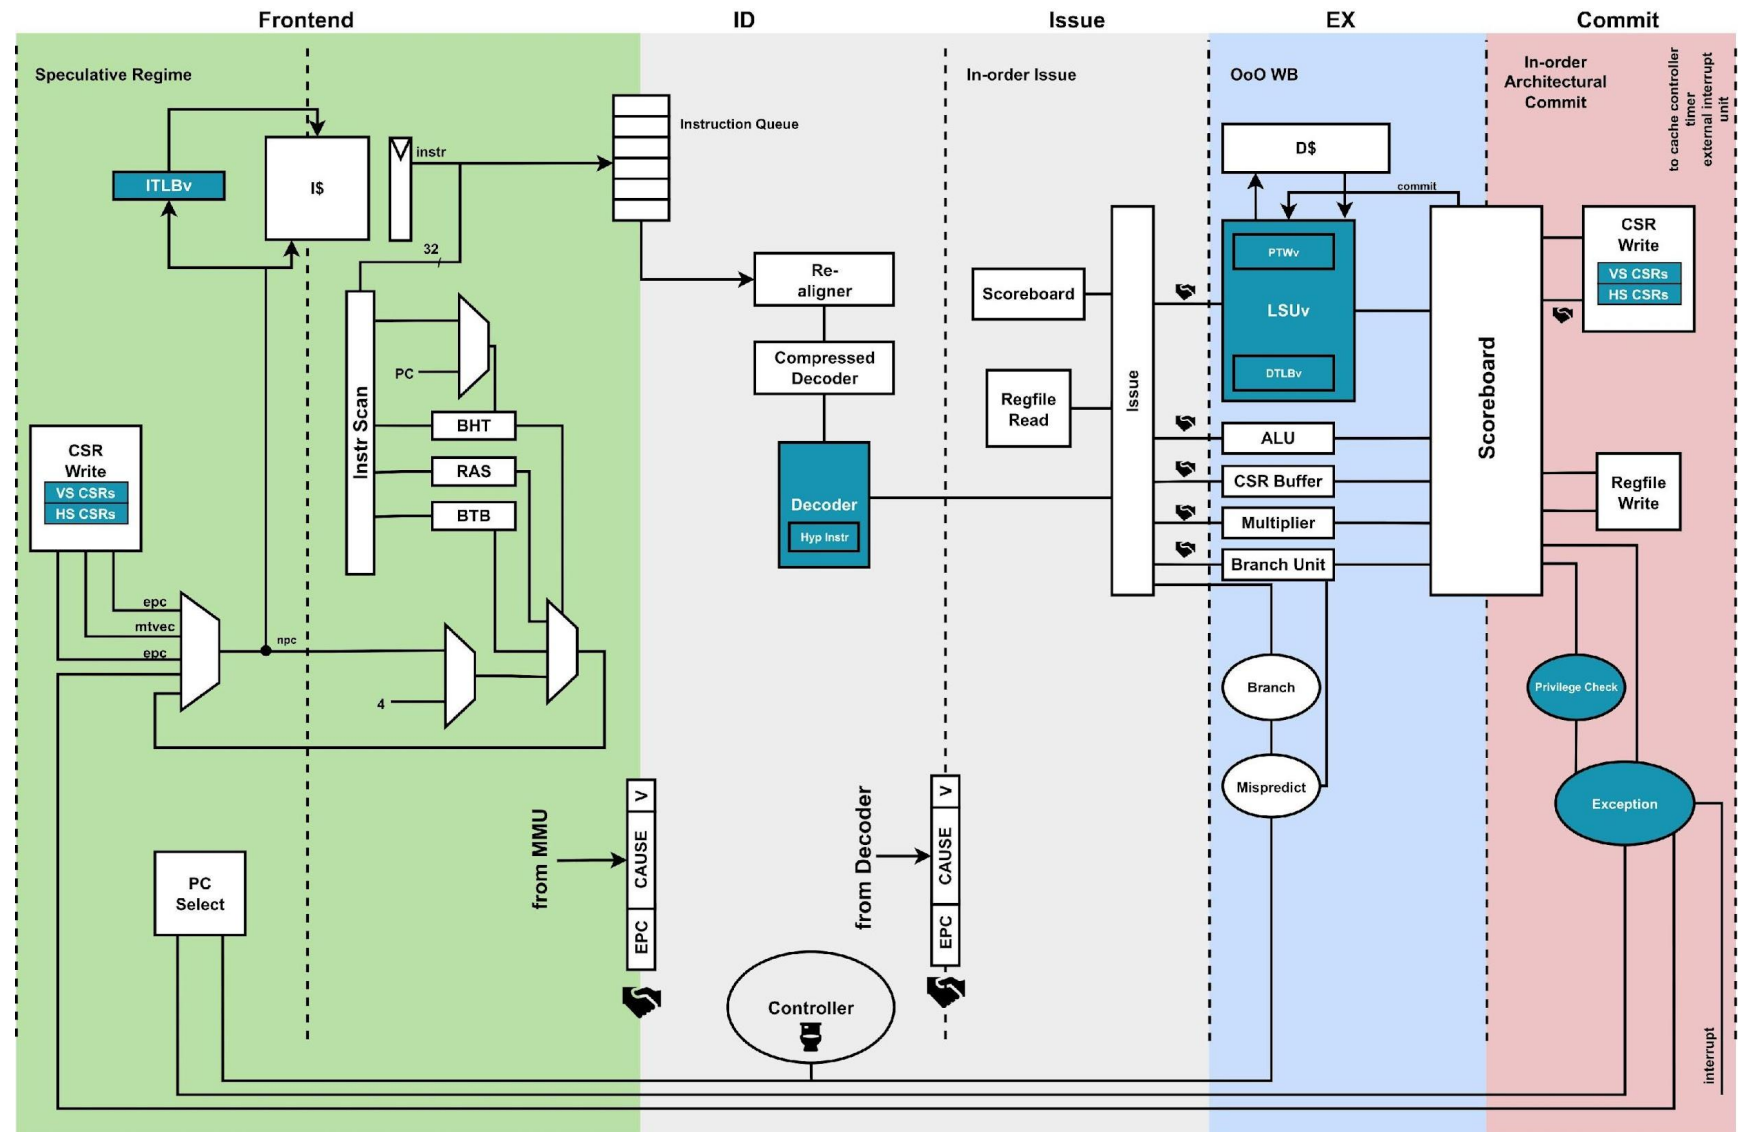
---

*The baseline privileged architecture is designed to simplify the use of classic virtualization techniques, where a guest OS is run at user-level, as the few privileged instructions can be easily detected and trapped. The hypervisor extension improves virtualization performance by reducing the frequency of these traps.*

*The hypervisor extension has been designed to be efficiently emulable on platforms that do not implement the extension, by running the hypervisor in S-mode and trapping into M-mode for hypervisor CSR accesses and to maintain shadow page tables. The majority of CSR accesses for type-2 hypervisors are valid S-mode accesses so need not be trapped. Hypervisors can support nested virtualization analogously.*

# CVA6 Hypervisor Extension

- CSR
- Decoder
- Exceptions
- PTW
- (I/D)TLB



# Specification Checklist

- H Extension, Version 1.0.0
- RV64 and sv39x4

Feature
Privilege Modes
Hypervisor and Virtual Supervisor CSRs
Hypervisor Instructions
Machine-level CSRs
Two-stage Address Translation
Traps

CSRs	hstatus/mstatus	●
	hideleg/hedeleg/mideleg	●
	hvip/hip/hie/mip/mie	●
	hgeip/hgeie	◐
	hcounteren	●
	htimedelta	●
	henvcfg	◐
	mtval2/htval	●
	mtinst/htinst	●
	hgapt	◐
	vsstatus/vsip/vsie/vstvec/vsscratch vsepc/vscause/vstval/vsatp	●
Instructions	hlv/hlvx/hsv	●
	hfence.vvma/gvma	●
Exceptions & Interrupts	Environment call from VS-mode	●
	Instruction/Load/Store guest-page fault	●
	Virtual instruction	●
	Virtual Supervisor sw/timer/external interrupts	●
	Supervisor guest external interrupt	●

Summary of market or input requirements

---

# Project Requirements

- Hypervisor Extension specification version 1.0
- Hypervisor Extension optional via parameterization
- Functional validated with multiple Hypervisors
- Functional validated with multiple Guest VM and Configurations:
  - Bare-metal Applications
  - Linux (CVA6-SDK)
  - FreeRTOS
- Tested and deployed in multiple FPGAs:
  - Genesys2
  - VCU118
- Extend CVA6-SDK (hypervisor)
  - Bao Hypervisor

# Potential Enhancements

## ■ CVA6 Hypervisor Extension RV32 Variant:

- RV32 Virtual Supervisor and Hypervisor CSRs
- MMU: SV32x4 translation scheme

## ■ MMU uArch Enhancements:

- PTE Cache
- PTW G-Stage TLB
- L2 TLB
- Svnapot

## ■ Relevant Extensions:

- Sstc (S- and VS-mode timer)
- Svnapot
- Svinval

## ■ Virtualization at the System Level:

- IOMMU
- AIA



# RISC-V AIA

## ■ Advanced Interrupt Architecture

- Advanced Platform-Level Interrupt Controller (APLIC)
- Incoming MSI Controller (IMSIK)
- Hardware assistance for VM interrupts
- Interprocessor Interrupts (IPIs)
- I/O MMU support for MSI to VMs

## ■ AIA spec

- Version 0.3.0 (AIA TG)
- Targeting ratification @ Q4 2022

## ■ Open question

- Interrupt virtualization support for APLIC

The RISC-V Advanced Interrupt Architecture


Document Version 0.3.0-draft

Editor: John Hauser  
jh.riscv@jhauser.us

June 13, 2022

# RISC-V IOMMU

- **Input-Output Memory Management Unit**
  - Support single or two-stage address translation
  - Supports VM and process isolation
  - Supports translation of AIA MSIs
  - Support for PCIe ATS and PRI
- **RISC-V IOMMU**
  - Version 0.1 -> not official spec (IOMMU TG)
  - Targeting ratification @ Q4 2022



# Project Impact and Industry Landscape

---

# Who would make use of OpenHW output ?

- Silicon Manufacturers
- Chip/SoC Designers
- FPGA Vendors
- Verification Companies
- Software Companies
- Universities / Research Centers


# RISC-V Hypervisors

Organization	Hypervisor	Status	License	Reference
Linux	KVM	Complete	GPLv2	GitHub: <a href="https://github.com/kvm-riscv/linux">https://github.com/kvm-riscv/linux</a> Tutorial: <a href="https://github.com/kvm-riscv/howto/wiki">https://github.com/kvm-riscv/howto/wiki</a>
Xvisor Project (Anup)	Xvisor	Complete	GPLv2	GitHub: <a href="https://github.com/avpatel/xvisor-next">https://github.com/avpatel/xvisor-next</a>
Bao Project / UMinho	Bao	Complete	GPLv2	GitHub: <a href="https://github.com/bao-project/bao-hypervisor">https://github.com/bao-project/bao-hypervisor</a>
Xen Project	Xen	WiP	GPLv2	GitHub: <a href="https://github.com/xen-project/xen">https://github.com/xen-project/xen</a>
Siemens	Jailhouse	WiP	GPLv2	GitHub: <a href="https://github.com/siemens/jailhouse">https://github.com/siemens/jailhouse</a>
seL4	seL4 VMM	PoC	BSD 2-Clause	Github: <a href="https://github.com/SEL4PROJ/sel4_riscv_vmm">https://github.com/SEL4PROJ/sel4_riscv_vmm</a>
RIVOS	Salus	WiP	-	Github: <a href="https://github.com/rivosinc/salus">https://github.com/rivosinc/salus</a>
fentISS	Xtratum NG	WiP	Commercial	Website: <a href="https://fentiss.com/products/hypervisor/">https://fentiss.com/products/hypervisor/</a>

# RISC-V Cores: Hypervisor Extension

Organization	Core	Status	License	Reference
CHIPS Alliance	Rocket	Complete	Apache-2.0	GitHub: <a href="https://github.com/chipsalliance/rocket-chip">https://github.com/chipsalliance/rocket-chip</a>
Cobham Gaisler	NOEL-V	Complete	Commercial / GPL	Website: <a href="https://gaisler.com/index.php/products/processors/noel-v">https://gaisler.com/index.php/products/processors/noel-v</a> Website (GRLIB): <a href="https://gaisler.com/index.php/downloads/leongrlib">https://gaisler.com/index.php/downloads/leongrlib</a>
SiFive	P200, P400, P500, P600	Complete	Commercial	Website: <a href="https://www.sifive.com/cores/performance">https://www.sifive.com/cores/performance</a>
StarFive	Dubhe	Complete	Commercial	Website: <a href="https://starfivetech.com/en/site/riscv-core-ip">https://starfivetech.com/en/site/riscv-core-ip</a>
InCore Semi	Chromite	WiP	BSD 3-Clause	GitLab: <a href="https://gitlab.com/incoresemi/core-generators/chromite">https://gitlab.com/incoresemi/core-generators/chromite</a>
IIT Madras	Shakti	WiP	BSD 3-Clause	GitLab: <a href="https://gitlab.com/shaktiproject/cores">https://gitlab.com/shaktiproject/cores</a>

# Why OpenHW Should Do This Project ?



# Project Members, Participants, and Timeline

---



# OpenHW Members Collaboration

- **Already committed:**

- University of Minho
- Zero-Day Labs

- **Possible collaboration:**

- Imperas (Verification?)
- UC Santa Barbara (Openpiton)
- Intel Pathfinder
- ( Others are welcomed :) )

# Project Team

- **Project Leader(s):**

- Sandro Pinto
- Bruno Sá
- José Martins

- **Technical Project Leader(s):**

- Sandro Pinto
- Bruno Sá
- José Martins

# Project Timeline



Q&A